

# CVIT Workshop Day 2

~~Session 1~~ session 1

## Document Images

skew correction

skew makes difficult to read & understand text

compute the skew angle/distortion angle, and then correct it

skew refers to the rotation of the document from its correct orientation

<https://tinyurl.com/24t6szta>

Colab Notebooks for Skew Correc

directly connect Colab to Github

create a folder named Day 2

the notebook would be updated on Github whenever you make a push.

descri

the libraries involved are skimage.transform

hough-line

hough-line-peaks

probabilistic-hough-line

rotate

canny

skimage.io.imread (can also use cv2, but it reads images in BGR)

skimage.color.rgb2gray

suggestion - read the documentations for the CV2 and skimage  
libraries

```
from google.colab.patches import cv2_imshow
```

## Steps

first mount the Google Colab Notebook to drive, and set up  
the environment to

### a code nuance

every  $\rightarrow$  image may not have just 3 channels

the given image just has 2 channels and if we use

$gray = \text{rgb2gray}(\text{image}[:, :, 3])$ , it would throw an  
error because the image just has 2 channels.

$\text{rgb2gray}$  only works for 3 color channels

~~so if there are only two color channels~~

there are usually four color channels, ~~there also~~ the fourth  
one being alpha channel.  $\text{rgb2gray}$  only works on RGB channels.

Hence we use  $(\text{this},)$  to use only the first 3 channels.

~~This~~ If the no. of channels is  $<= 2$ , then the  $gray$  is simply  
the input image itself. Thus we would have

If  $\text{len}(\text{image}, \text{shape}) > 2$ :

$$\text{gray} = \text{rgb2gray}(\text{image}[:, :, :3])$$

else:  $\text{gray} = \text{image}$

## Edge Detection - Canny

remember Sobel filters?

Canny Edge Detection is similar with extra preprocessing and refinement steps like noise reduction, non-maximal suppression, etc.

If the edges are very fine, very sharp then we reduce the size of the filter, and vice versa.

Function signature: canny(image, sigma=1.0, low\_threshold=10,  
high\_threshold=None, mask=None, use\_quantiles=False,  
(mode='constant', cval=0.0))

## Canny Edge Detection - OpenCV Docs

multi-stage algorithm

1) Noise Reduction

2) Finding intensity gradient of the image

$$\text{Edge-gradient } (G) = \sqrt{G_x^2 + G_y^2}$$

$$\text{Angle } (\theta) = \tan^{-1} \left( \frac{G_y}{G_x} \right)$$

The notebook is written in skimage, but mostly CV2 is used for most purposes. The canny can also be implemented in CV2 instead of skimage.

③ non-maximum suppression

④ hysteresis thresholding

### Noise Reduction

remove noise with a  $5 \times 5$  Gaussian filter

### Intensity Gradient

non maximum suppression means suppressing values other than the maximum values. So suppose we have 3 values, ~~or~~ high values. ~~so~~ so we might select the highest one and neglect the rest.

### Hough-Transform

details ~~will~~ will be dealt with in DIP course.

used to detect geometrical shapes — circles, triangles

how? parametric space ~~and~~ ~~will~~ based on accumulator

in any kind of transform first we define it mathematically  
for e.g. circle needs three values —  $(x, y, R)$   
 $\downarrow$   
centre coordinates

for a line we need a slope and a point ~~or~~

"image processing origin lies on the top-left corner"

to detect lines, we calculate the edges. we take a point, and see if a line passes through the point

np.linspace is used

angles = np.deg2rad(np.arange(0.1, 180.0))

for any point, if a lot of lines are passing through it, then it is a point of a structure

h-accumulator

hough space is  $(r, \theta)$

every  $(r, \theta)$  is a line, which might be passing through a lot of pts. say 60, then h is 60, because the line got 60

transform to Hough space to because in the Cartesian space the complexity to identify line would be  $180 \times (\text{no. of points})^2$

If the image size increase, the search space for the Hough space does not increase.

Hence computation becomes much easier.

h.shape(2001, 180)

distance from origin

h is a matrix of this size. Each value corresponds to the no. of points having that particular  $(r, \theta)$ .

We want to know those parts with the maximum points.

We had given the resolution for theta (during & lenspace).  
 but we did not give the resolution for r.  
each matrix value  
 how many points lie on that line

### hough-line-peaks

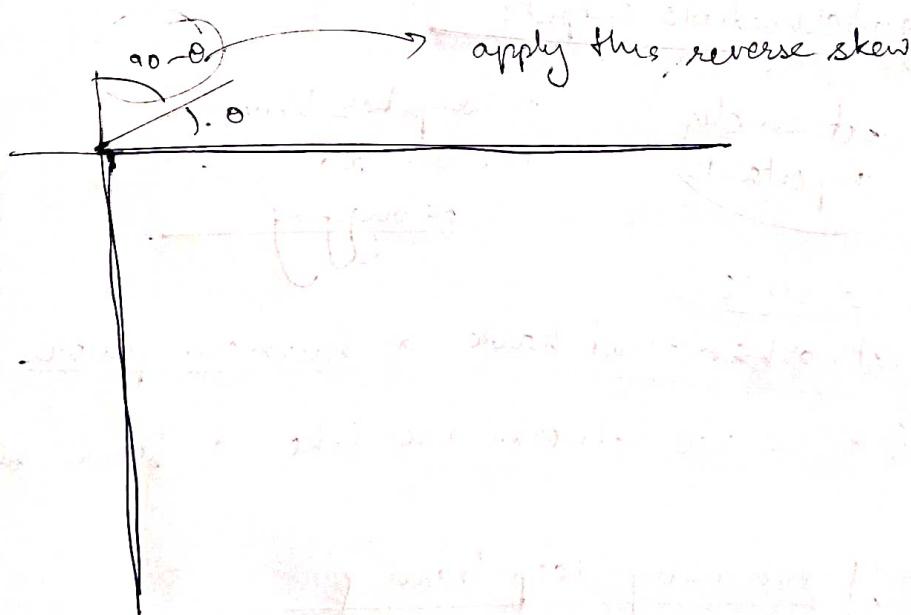
takes as input the hough-line-space

accum, angles, dists & hough-line-peaks ( $h, \theta, d$ )

list ( $\text{zpl}^* \text{hough-line-peaks } (h, \theta, d)$ )

still there are outliers.

use the property that maximum lines obtained would be correctly oriented.



affine transformation = translation + rotation + scaling

### translate

$$\begin{pmatrix} x' \\ y' \\ 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

### 2D rotate

$$2D \begin{pmatrix} x' \\ y' \\ 1 \end{pmatrix} = \begin{pmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

shear

$$\begin{pmatrix} x' \\ y' \\ 1 \end{pmatrix} = \begin{pmatrix} 1 & sh_x & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

rotation is a combination of two shears

cv2.getRotationMatrix2D (center, skew-angle, 1.0)

center + interpolation done in the ~~image~~ while rotation

show → return → false

shows all intermediate steps for debugging.

visualizing intermediate outputs

(extremely important)

in Computer Vision  
one can do visual\*  
debugging.

General advice

neural networks do not have a human process

overall neural networks are like a black ~~box~~ box

disadvantages of not using deep learning

- too powerful.
- compute intensive
- huge amounts of data

## second notebook

<https://tinyurl/bddrcygt>

### OCR

word level segmentation } both done in grayscale only  
line level segmentation }

done in CV2.

### Steps

gray = cv2.cvtColor (image, cv2.COLOR\_BGR2GRAY)

def binarize\_and\_word\_detect\_contours

# load & convert to grayscale

# denoise the image

1) denoise ~~the~~ with Gaussian blur, kernel 5x5

2) advanced denoising technique

cv2.fastNlMeansDenoising (blurred, None, h=10,  
templateWindowSize

denoising necessary to segregate words, because word identification requires detection of a continuous white space.

If we did not denoise, then multiple words could be detected as one.

# thresholding

cv2.threshold (denoised, 127, 255, cv2.THRESH\_BINARY)

# make the text thicker

we apply a dilation operation

there are white spaces within ambiguous text, to remove these we use dilation operation

dilated-image = cv2.dilate (binary-image, text-kernel, iterations)

return dilated-image

function

getStructuringElement defines a kernel

cv2.getStructuringElement (cv2.MORPH\_RECT, (H, W))

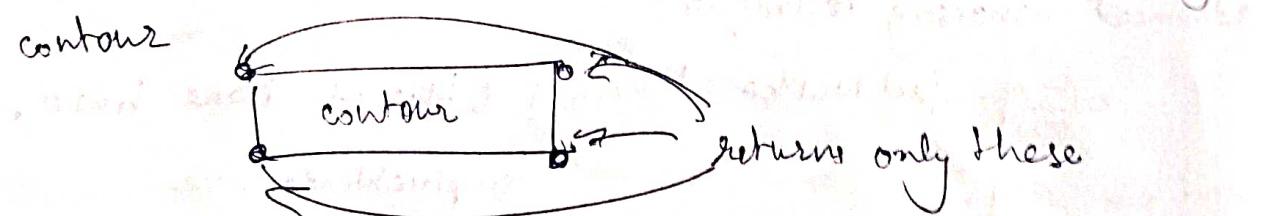
We will use a rectangular kernel here, to because this is being done for words which are rectangular roughly

contours = cv2.findContours (wordTextDocument, cv2.RETR\_EXTERNAL  
cv2.CHAIN\_APPROX\_SIMPLE)

(What is returned?)

RETR\_EXTERNAL retrieves only the external points of each

contour



CHAIN\_APPROX\_SIMPLE: only the four corners are computed

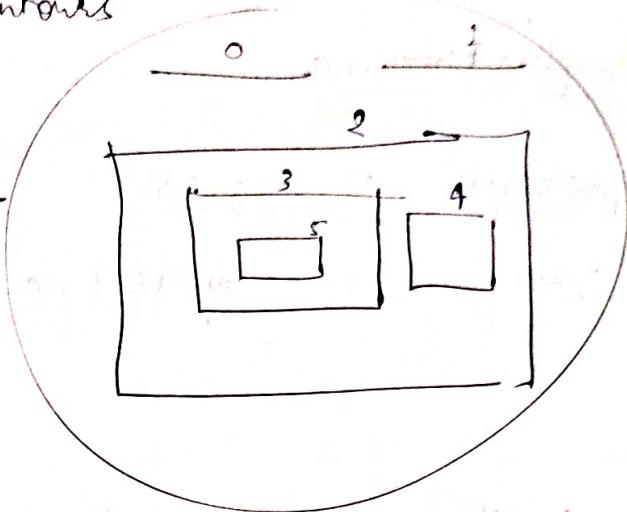
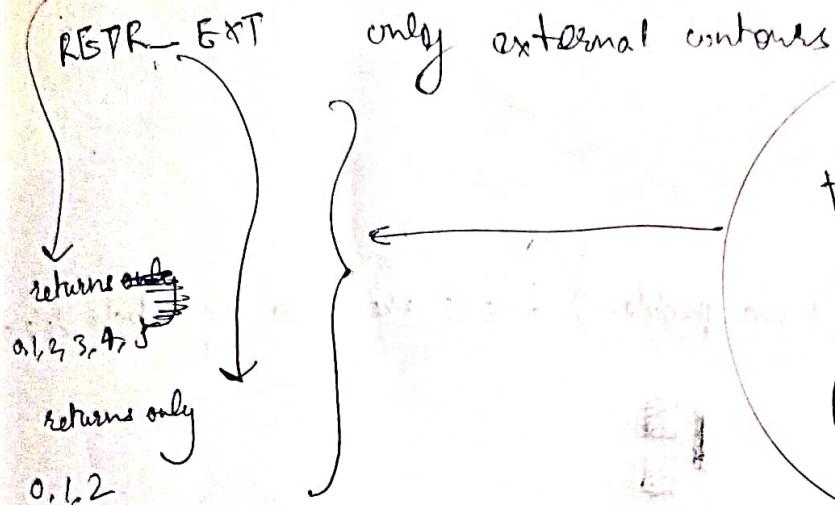
(What is computed?)

drawContours

(draw these contours)

Contours Hierarchy

RETR\_LIST returns with no hierarchy, but all ~~internal~~ internal contours also returned



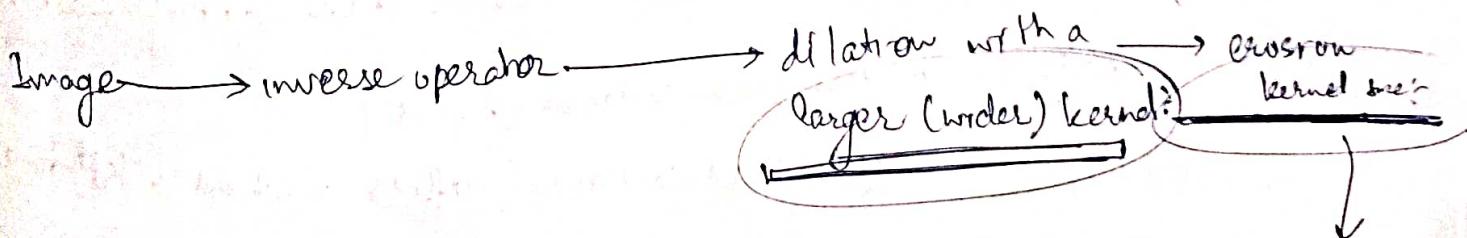
now we have got all the words segmented from each other.

Suppose we wanted EXACT boundaries of the ~~the~~ words then after having treated these word blobs. After dilation we dilate with a bigger kernel to get only the word blob and then do fine contouring.

Now, we can pass each word separately into the deep learning model.' word level OCT

for line segmentation we would use a kernel with a larger width.

Then,



Erosion & dilation can be used as noise removal techniques also.

<https://tryuel.com/st9ptrmp>

CVIT Workshop Day 2

Session 2

extracting - a

## Template Matching

problem with cheques

apart from the ~~no~~ digits (no problem) there are also symbols like



which is a single symbol.

partwise matching — break an image into smaller parts, template match each of these individual parts

we will first have a reference image containing all characters & digits in the correct order. This will be a clear image

### Steps

1) extracting the bounding boxes from the reference image

cv2.findContours returns two values, the contours & the hierarchy

~~or~~ we can

refCnts: cv2.findContours(`ref.copy()`, cv2.RETR\_EXTERNAL)

cv2.CHAIN\_APPROX\_SIMPLE

refCnts: imutils.grab\_contours(`refCnts`)

refCnts: contours.sort\_contours(`refCnts`, method = "left-to-right")

def extract\_digits\_and\_symbols:

iterate over the contours.

rois = [ ]

regions of interest

the region inside the bounding box

the four corners are enough to get the region of interest.

two other params:

min height, min width

use the fact that the height & width of the "other" character symbols are much smaller than the digits.

On each contour, we use

$(cx, cy, cw, ch) = cv2.boundingRect(c)$

to obtain the

if  $cw \geq mnW$  and  $ch \geq mnH$ :

# it is a digit

roi = rimage[cy:cy+ch, cx:cx+cw]

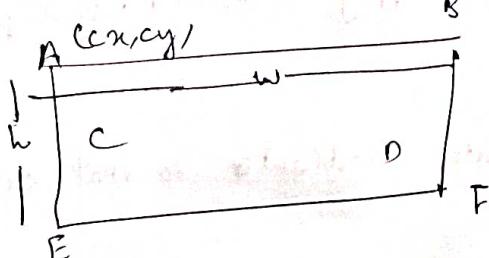
rois.append(roi)

} gives the origin coordinates,  
the height, the width of the  
bounding box of the  
contours

Note that we had resized the image in the beginning so that we can set the hyperparameters such as minW & minH.

More detailed explanation

suppose we have 6 contour points:



we would be given returned  
 $cx, cy, cw, ch$  using  
 $cv2.boundingRect$ .

Now, once we have segregated the digits, we need to segregate the characters which have more than one continuous regions.

Note that each of them is just 3 regions large. So just combine 3 at a time when you come to characters.

Now we map the sets to the names of the digits/characters represented by them.

Now, we again perform a resize so that template matching is possible.

Now coming to the Query Image

We will first crop the bottom part of the cheque because that contains the ~~main~~ code.

first convert to grayscale

then perform ~~morphological operations~~ morphology operation:

blackhat = cv2.morphologyEx (gray, cv2.MORPH\_BLACKHAT, kernel)

finds the dark regions against a light.

Note that we could also have performed a binary inversion, but morphology operation is better in most cases.

~~# compute the~~

Now we need to extract the bounding boxes of the query image.

First we can divide the whole cheque no. into three words/group

~~for~~ So we would need to apply dilation (usual procedure)

BUT

first we apply a horizontal gradient, so that all the horizontal

So we apply a Sobel filter to remove external edges, which includes watermark  
the highest values obtained by this are made 255, others are made 0.

Closing (dilation followed by erosion)

closing small holes in the foreground images

cv2.morphologyEx(gradient, cv2.MORPH\_CLOSE, rectKernel)

thresholding

move the whites to 255, rest to 0

skimage.clear\_border

this removes the "external borders" in the image

does not remove the lower boundary because it is not continuous,  
thus it wants to detect a continuous boundary

(Note that these operations are best performed on binary images with only  
0 & 1)

To separate the rest of the contours from those corresponding to the  
characters/digits, we use the width/height thresholding

(the contours corresponding to the edges have much smaller contour size  
than the characters/digits)

Now we need to get the context inside the bounding boxes around the  
bundles of characters/digits that we just formed

We again draw contours and bounding boxes around each character.

We will now template match with our stored information from the reference image.

result = cv2.matchTemplate

cv2.TM\_CCOEF : correlation coefficient

cv2.TM\_CCORR : cross correlation



$$RC(x, y) = \sum_{x', y'} (f(x', y') \cdot g(x+x', y+y'))$$

normalized by the size of the kernel

$$RC(x, y) = \sum_{x', y'} (\underbrace{f(x', y')}_{\text{template}} \cdot \underbrace{g(x+x', y+y')}_{\text{Image}})$$

output of ~~match~~ template matching will be another image.

If we are doing difference matching, then a black dot / region in the matching image would correspond to a good match.

minLoc, maxLoc

cv2.minMaxLoc(result)

and minMaxLoc(result)

however template matching is not ideal for real life images matching.

In these cases it would be best to train your own Holt Linear SVM classifier or a CNN.

## Prof. Ran Kiran S Presentation

less availability of data

aspect ratio extremely high

need a method which does not require resizing.

because manuscripts come in different sizes

just to segment the lines

### solutions

1) increase ~~max~~ each image to max width & max height

problem due to huge compute required

image grainy problem

2) breaking horizontally (patching)

### seen understanding

find all minimum energy path calculation and remove them from the image

take away :- the mindset, the thought process  
and not quite so much the technique